

# Computer Networks

## Chapter 2.5

Jong-won Lee

Handong University

# Chapter 2: Application layer

- ❑ 2.1 Principles of network applications
- ❑ 2.2 Web and HTTP
- ❑ 2.3 Electronic Mail (SMTP, POP3, IMAP)
- ❑ 2.4 DNS
- ❑ 2.5 P2P applications
- ❑ 2.6 Video streaming and content distribution networks (CDNs)
- ❑ 2.7 Socket programming with TCP and UDP

# Properties

- ❑ No central control, no central database
- ❑ No hierarchy
  - Every node is both a client and a server
  - The communication between peers is symmetric
- ❑ No global view of the system
  - Scalability
- ❑ Availability for any peer
- ❑ Peers are autonomous
- ❑ System globally unreliable
  - Robustness and security issues

# Napster: centralized directory


original "Napster" design

1) when peer connects, it informs central server:

- IP address
- content

2) Alice queries for "Hey Jude"

3) Alice requests file from Bob




# Napster: centralized directory

- ❑ File-sharing system
- ❑ Almost distributed system
  - The location of a document is centralized
  - The "transfer" is peer-to-peer
  - Fast querying
- ❑ Problems
  - Robustness
    - Single point of failure
  - Scalability (?)
    - Performance bottleneck
  - Copyright infringement
    - Sentenced to go out of business


# Gnutella: Query flooding

- ❑ fully distributed P2P protocol
  - no central server
- ❑ public domain protocol
  - many Gnutella clients implementing protocol
- ❑ Gnutella protocol
  - Based on broadcast/back-propagation mechanism over an overlay network
  - Message types
    - Ping/Pong: for group membership
    - Query/Query Hit: for search

# Overlay Networks



# Overlay Networks

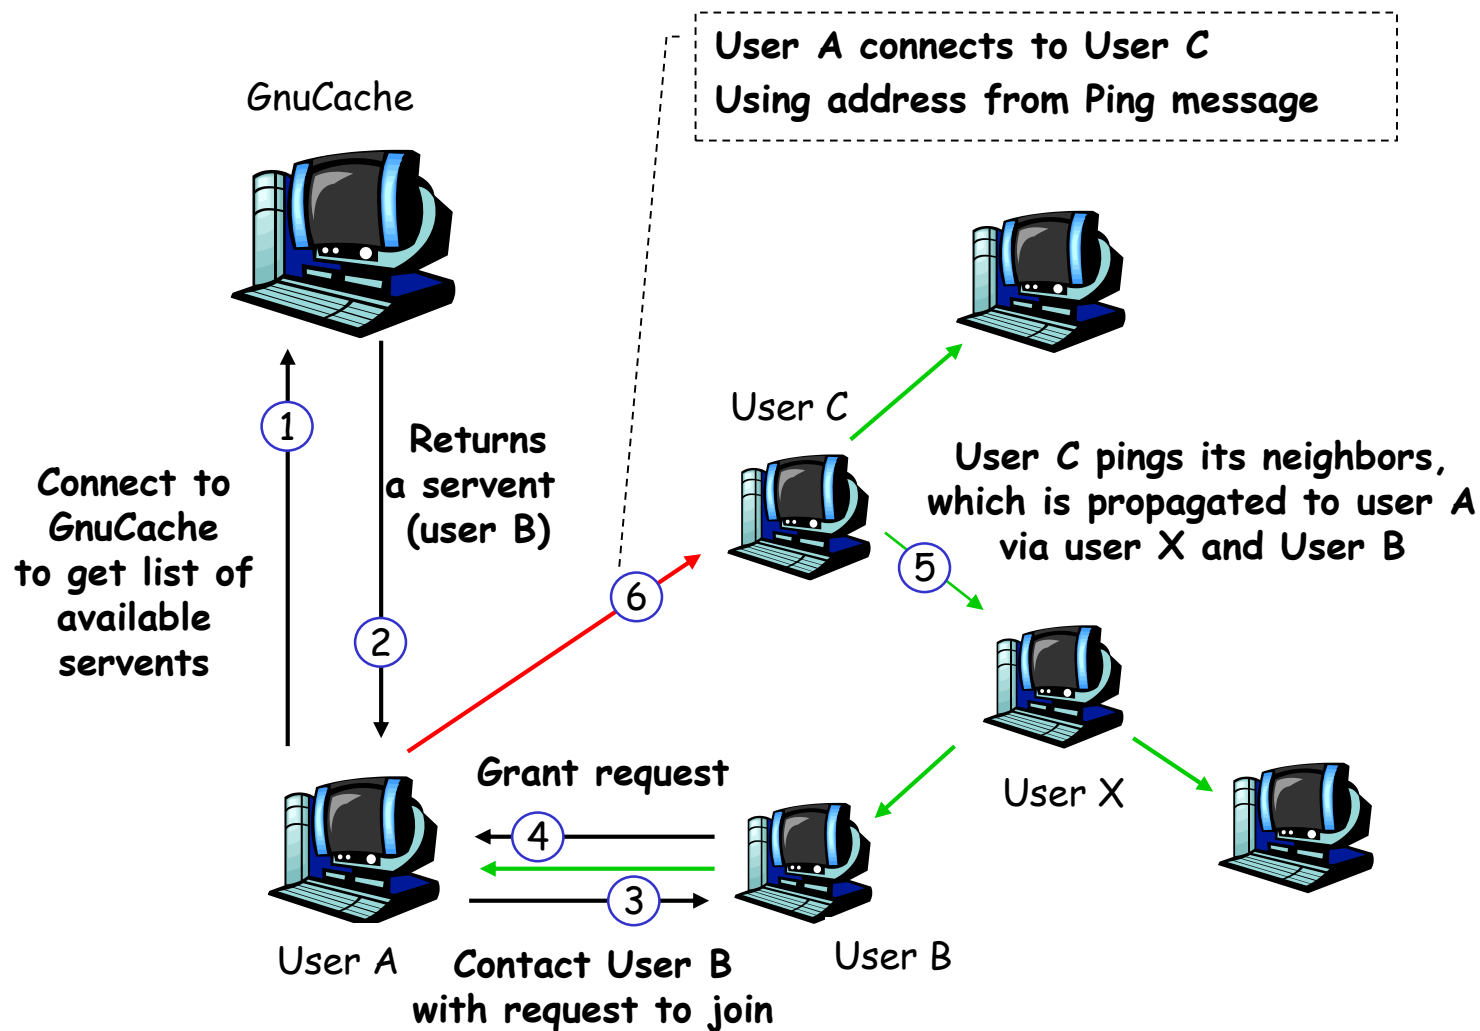




# Gnutella: Join operation

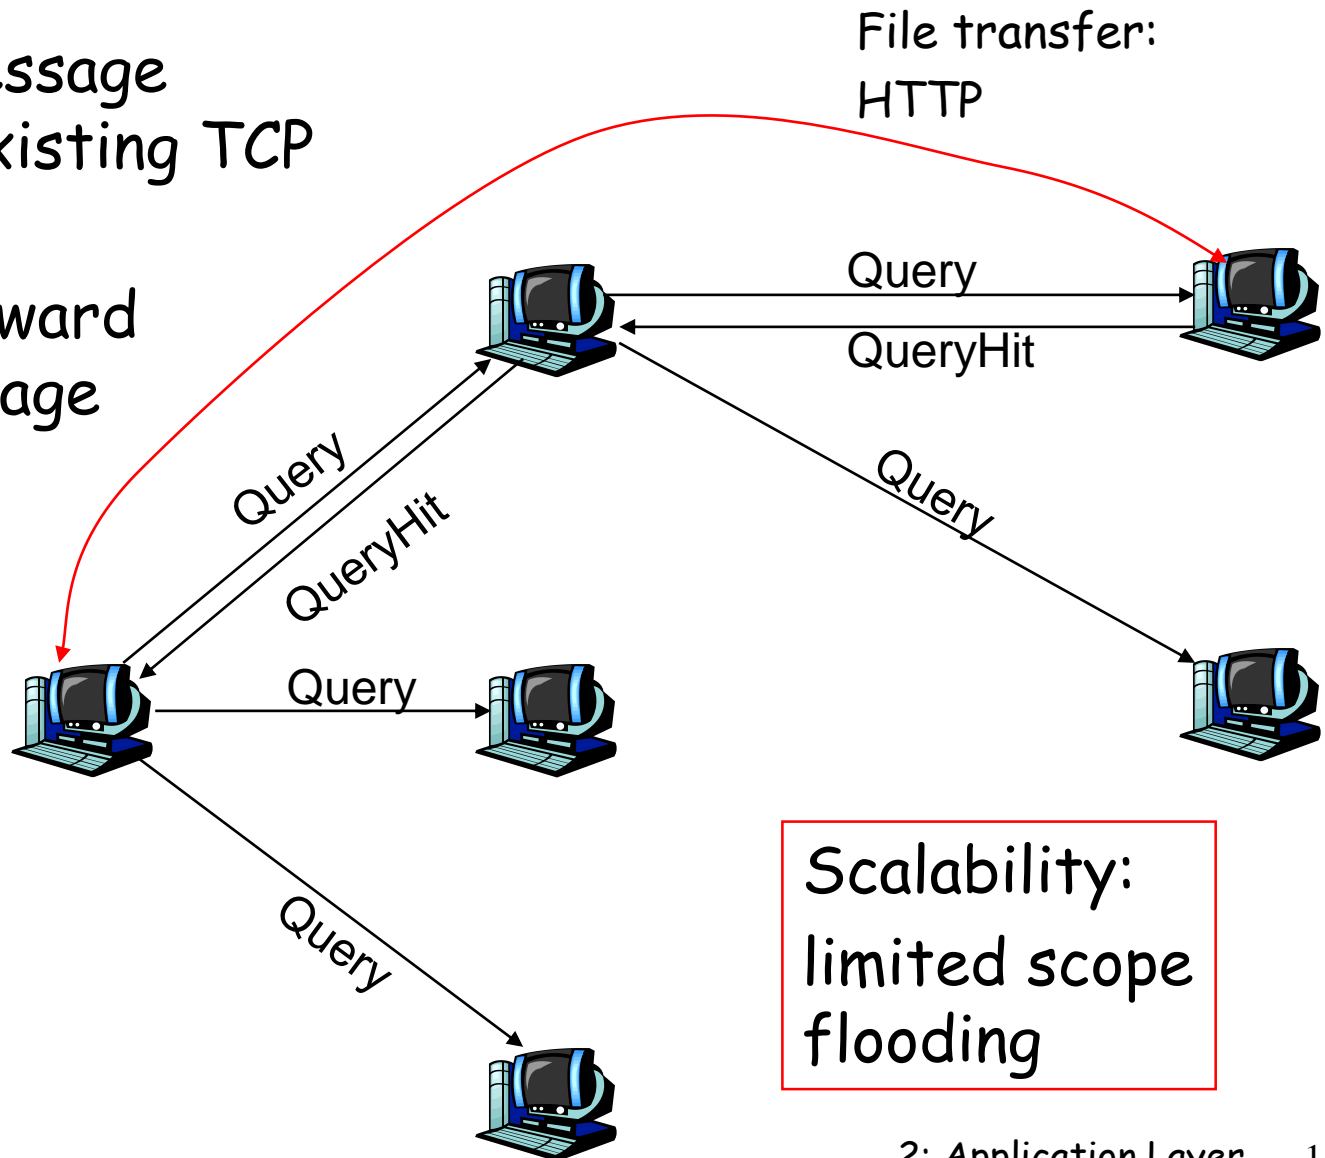
1. Joining User A must find some other peer in Gnutella network: use list of candidate peers. In this case, the user connects to a *GnuCache* server.
2. *GnuCache* returns a list of nodes on the Gnutella network. User A chooses one of these (User B) and attempts to contact it
3. User A sends a "Gnutella Connect" (see section 4) to User B to request to join the Gnutella network
4. User B accepts and returns a "Gnutella OK" to User A. User A is now part of the Gnutella network and is connected to one other Gnutella node.
5. Peers in Gnutella *Ping* their neighbors with their information periodically. Such *Ping* messages are not only replied to with *Pong* messages but they are propagated along to all other interconnected *servents*.
6. In this way, User A finds out about User C because User C has propagated its *Ping* message over to User A. The *Ping* messages contain the address of the sender. Typically Gnutella *servents* connect to around 3 other *servents* in the network.

# Gnutella: Join operation



# Gnutella: Search

- ❑ Query message sent over existing TCP connections
- ❑ peers forward Query message
- ❑ QueryHit sent over reverse path




# Gnutella

- ❑ Gnutella is simple, however
- ❑
- ❑ Excessive query traffic
  - Can be controlled through (limited-scope query flooding)
- ❑
  - Limited flooding reduces the number of peers
- ❑ Maintenance of overlay network
  - TCP connection between peer neighbors should be maintained even when there is no traffic
- ❑ Free riding
  - Most Gnutella users do not provide files to share
  - 47% of all responses are returned by top 1% of hosts

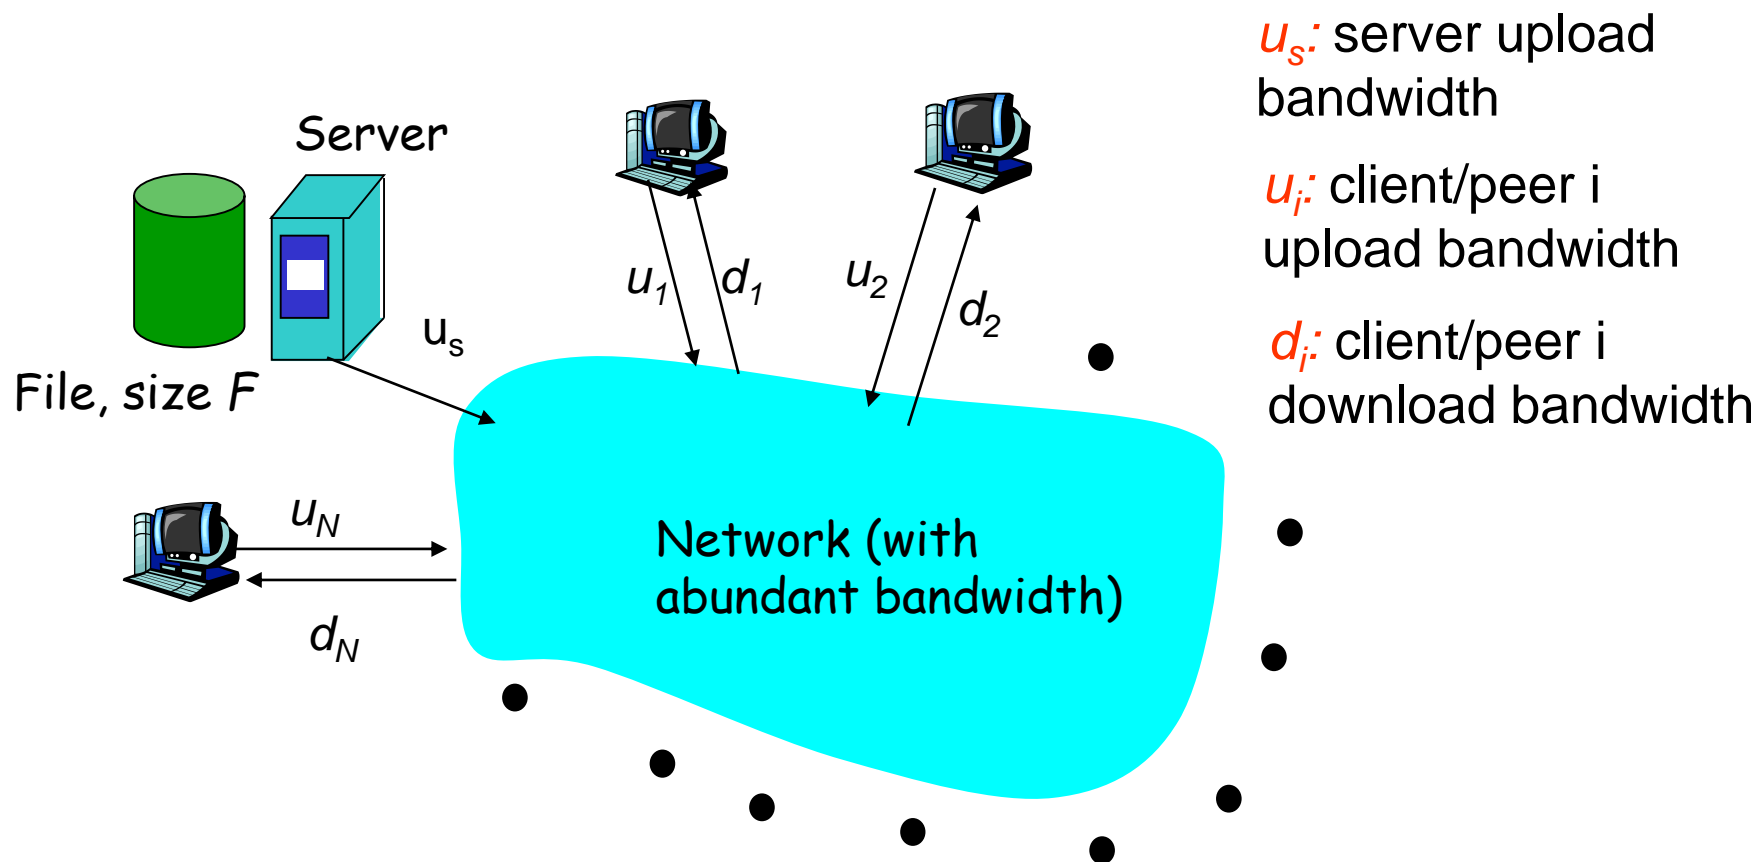
# Hierarchical Overlay

- ❑ between centralized index, query flooding approaches
- ❑ each peer is either a *group leader (super-peer)* or assigned to a group leader.
  - TCP connection between peer and its group leader.
  - TCP connections between some pairs of group leaders.
- ❑ group leader tracks content in its children




# Comparing Client-server, P2P architectures

Question : How much time to distribute file initially at one server to  $N$  other computers?



# Client-server: file distribution time

- server sequentially sends  $N$  copies:
  - $NF/u_s$  time
- client  $i$  takes  $F/d_i$  time to download




Time to distribute  $F$   
to  $N$  clients using client/server approach  
 $= d_{cs} = \max \{ NF/u_s, F/\min_i(d_i) \}$

increases linearly in  $N$   
(for large  $N$ )

# P2P: file distribution time


- ❑ server must send one copy:  $F/u_s$  time
- ❑ client  $i$  takes  $F/d_i$  time to download
- ❑  $NF$  bits must be downloaded (aggregate)
  - ❑ fastest possible upload rate (assuming all nodes sending file chunks to same peer):  $u_s + \sum_{i=1,N} u_i$



$$d_{\text{P2P}} \geq \max \left\{ F/u_s, F/\min(d_i), NF/(u_s + \sum_{i=1,N} u_i) \right\}$$



# Comparing Client-server, P2P architectures




# P2P Case Study: BitTorrent


## □ P2P file distribution

tracker: tracks peers participating in torrent

torrent: group of peers exchanging chunks of a file



# BitTorrent (1)



- ❑ file divided into 256KB *chunks*.
- ❑ peer joining torrent:
  - has no chunks, but will accumulate them over time
  - registers with tracker to get list of peers, connects to subset of peers (“neighbors”)
- ❑ while downloading, peer uploads chunks to other peers.
- ❑ peers may come and go
- ❑ once peer has entire file, it may (selfishly) leave or (altruistically) remain

# BitTorrent (2)

## Requesting chunks


- ❑ at any given time, different peers have different subsets of file chunks
- ❑ periodically, Alice asks each neighbor for list of chunks that they have.
- ❑ Alice issues requests for her missing chunks
  -

## Sending Chunks:

- ❑ Alice sends chunks to **four neighbors** currently **sending her chunks *at the highest rate***
  - ❖ re-evaluate every 10 secs
- ❑ every 30 secs:
  - , starts sending chunks
  - ❖ newly chosen peer may join top 4

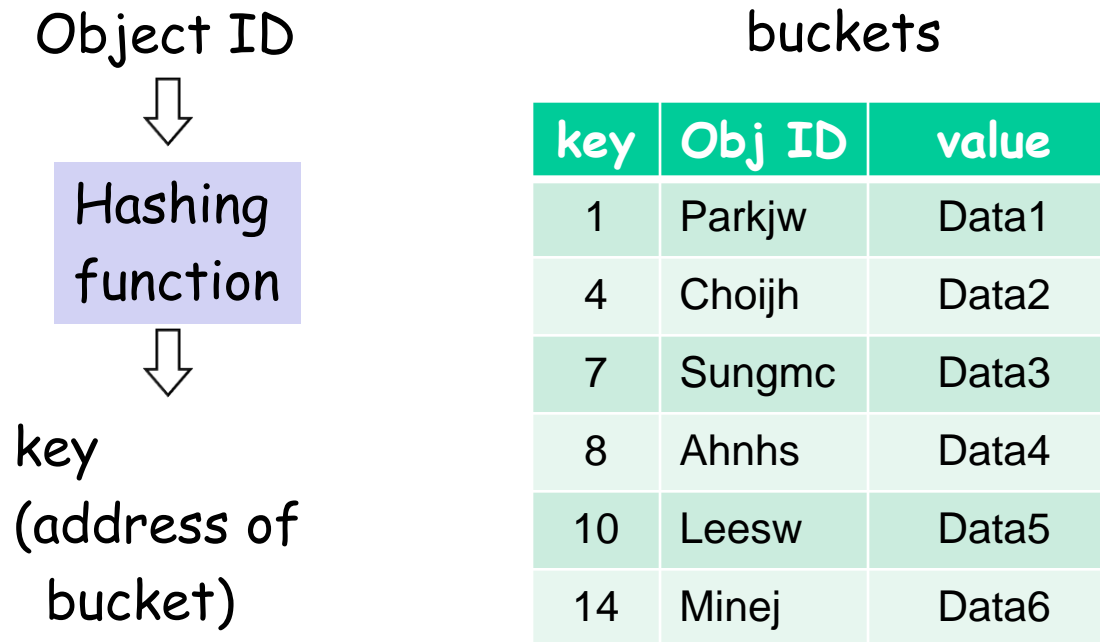
# BitTorrent: Tit-for-tat

- (1) Alice “optimistically unchokes” Bob
- (2) Alice becomes one of Bob’s top-four providers; Bob reciprocates
- (3) Bob becomes one of Alice’s top-four providers



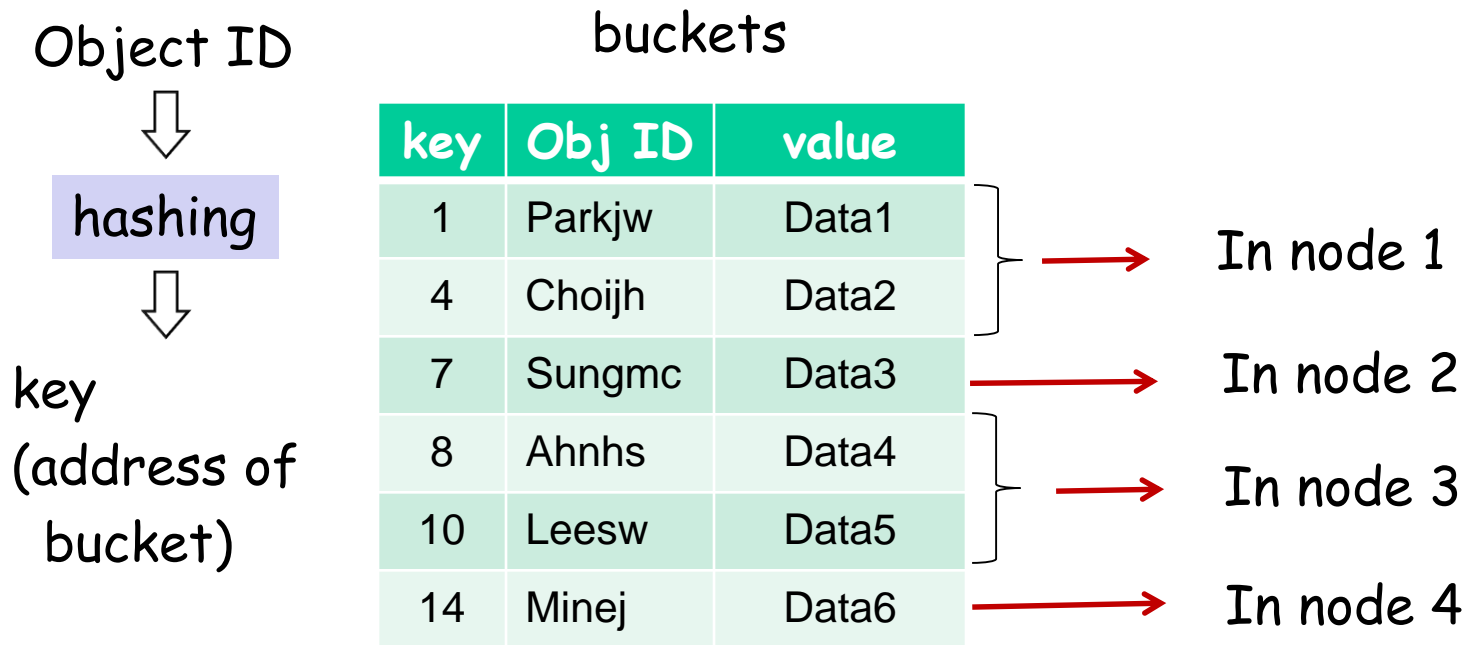
# Distributed Hash Table (DHT)

- An ordinary hashtable, which is ...



# Distributed Hash Table (DHT)

- An ordinary hashtable, which is **distributed**



# Distributed Hash Table (DHT)

- ❑ DHT = a distributed P2P database
  - Distributes data among a set of nodes according to predefined rules
- ❑ Database has (object ID, value) pairs;
  - Object ID: ss number; value: human name
  - Object ID: movie title; value: IP address
- ❑ Peers query DB with object ID or key
  - DB returns values that match the key
  - In DHT-based networks each peer has a partial knowledge about the whole network. This knowledge can be used to route the queries to the responsible nodes using effective and scalable procedures.



# Q: how to assign keys to peers?

## ❖ central issue:


- assigning (key, value) pairs to peers.

## ❖ basic idea:

- convert each key to an integer
- Assign integer to each peer
- put (key,value) pair in the peer that is **closest** to the key


# DHT

- Peer IDs and object IDs are in the same range
  - Assign an integer to each peer or each object in range  $[0, 2^n - 1]$ .
  - $\text{key} = \text{hash}(\text{object name})$




Set of nodes

Key of nodes



Set of items

Key of items



1. Decides on common key space for nodes and items

2. Connects the nodes smartly

3. Make a strategy for assigning items to nodes

# DHT

## □ Consistent hashing

- A scheme that provides hashing table functionality in a way that the addition or removal of one node does not significantly change the mapping of keys to nodes

# Consistent hashing using a ring

- The keys of Peer(node) IDs and object IDs are in the same range
  - Ex.: the size of key space = 16 ( $[0, 15]$ )

csee.handong.edu



$H(\text{csee.handong.edu})$   
=12

salt.handong.edu



$H(\text{salt.handong.edu})$   
=3

light.handong.edu




$H(\text{light.handong.edu})$   
=0

203.252.100.60




$H(203.252.100.60)$   
=7

comnet\_lec1.pdf



$H(\text{comnet\_lec1.pdf})$   
=2

movie1.avi



$H(\text{movie1.avi})$   
=12

music1.mp3




$H(\text{music1.mp3})$   
=14

# Consistent hashing using a ring


## □ Node connections

- The successor node of node  $i$  is the first node with ID greater than  $i$ .
- The ID range is considered as a circular space.



# Consistent hashing using a ring

- The policy that each item is assigned to a node
  - An item with ID  $x$  is assigned at the node with ID  $\text{succ}(x)$ .
  - Def.  $\text{succ}(x)$  is the first node on the ring with ID greater than or equal to  $x$  in the fashion of circular space.



# Consistent hashing using a ring

- ❑ If each node knows its successor node, the two operations, `get()` and `put()`, would be simply done by searching them sequentially.
  - ❑ `Put(hash(item),value), get(item)`
- ❑ The average number of messages to resolve a query
  - $O(N)$  :  $N$  = the number of nodes
  - Possible to reduce to  $O(\log N)$ 
    - How? (refer to Chord scheme)
- ❑ Consider how to handle peer